

HUMAN vs AI IMAGE DETECTION

EdX Data Visualization &
Analytics

MARCH 2025

Team Members

Jacob Sickerman

Lannon Grady

Matthew Matti

Maxwell Leuthner

Toka Hassan

Toni Makakoa

Process

Download & Load Kaggle Data

Create, Train, & Test Learning Model

Improve the Learning Model

Create Flask App to Upload Image & Predict

Web Page Deployment

AI vs. Human-Generated Images

A Curated Dataset of AI-Generated and Authentic Images



Data Card Code (92) Discussion (0) Suggestions (1)

About Dataset

Official dataset for the 2025 Women in AI Kaggle Competition:

<https://www.kaggle.com/competitions/detect-ai-vs-human-generated-images>

The dataset consists of authentic images sampled from the Shutterstock platform across various categories, including a balanced selection where one-third of the images feature humans. These authentic images are paired with their equivalents generated using state-of-the-art generative models. This structured pairing enables a direct comparison between real and AI-generated content, providing a robust foundation for developing and evaluating image authenticity detection systems.

Examples of Images in Dataset

```
# Show examples
import matplotlib.image as mpimg

def image_examples(image, num_images=5):
    for j in range(2):
        example_image_path = f"./Resources/TrainingImages/{j}/"
        fig = plt.figure(figsize=(15, 5))
        for i in range(5):
            ax = fig.add_subplot(1, 5, i+1)
            img = mpimg.imread(f'{example_image_path}/{os.listdir(example_image_path)[i]}')
            ax.imshow(img)
            ax.axis("off")
            plt.title("Human-made" if j == 0
                      else "AI-generated")
        plt.show()
image_examples(image)
```

The dataset on Kaggle
contained 79,950
images.

50% human-created

50% AI-generated

Examples of Images in Dataset

Different sizes,
different aspect
ratios

Human-made



Human-made



Human-made



Human-made



Human-made



AI-generated



AI-generated



AI-generated



AI-generated



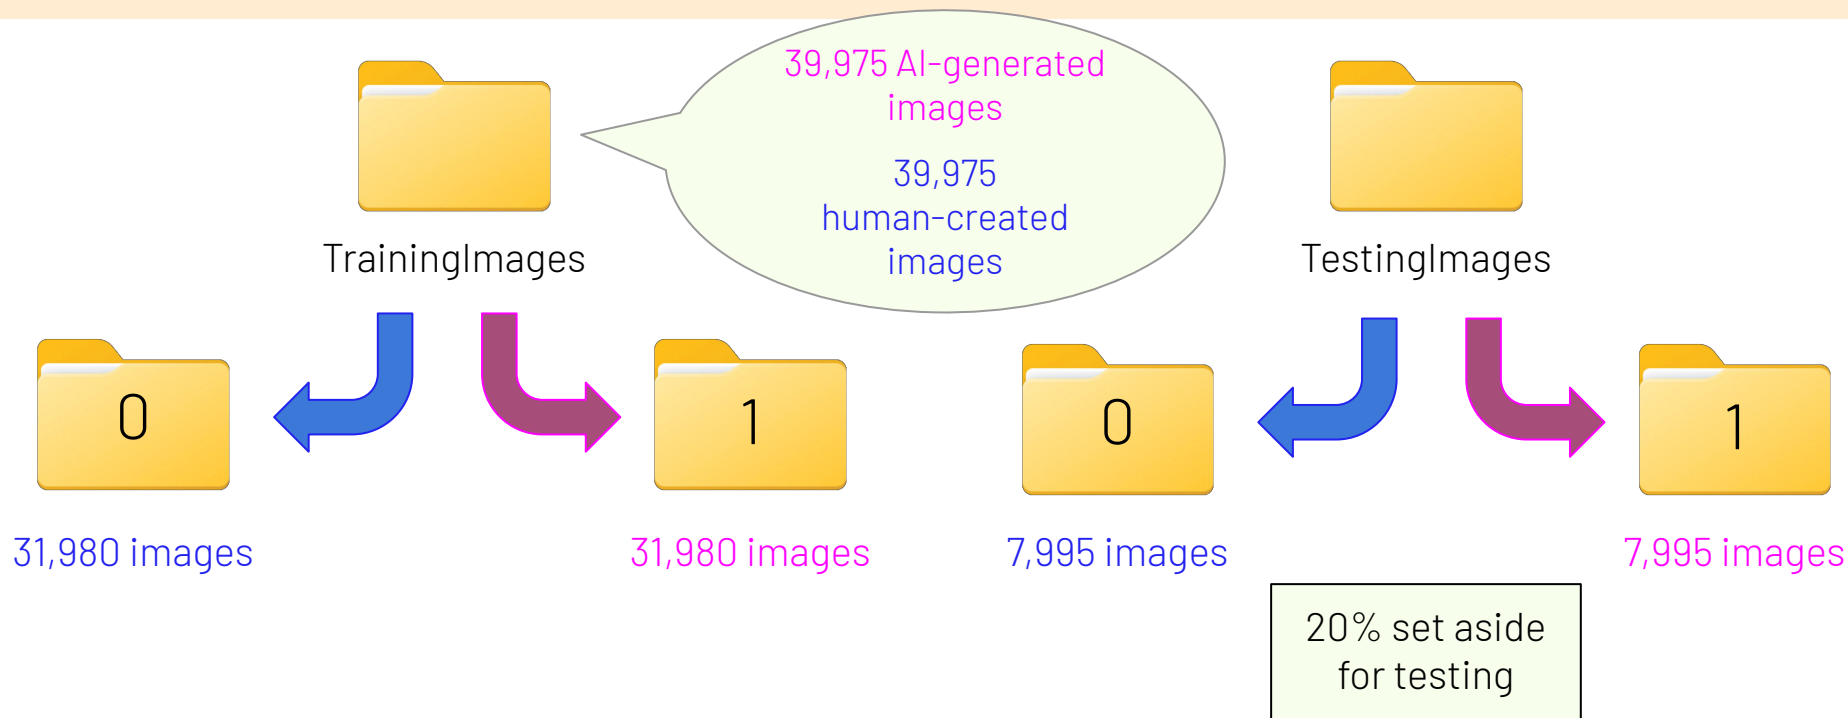
AI-generated



Creating, Training, & Testing the Learning Model

We chose to use a Convolutional Neural Network because they are extremely successful at image processing, recognition, and classification, making it the best choice to help us predict whether an image is AI vs. Human-Generated.

Splitting Testing & Training Data



Splitting Testing & Training Data

```
import os
import pandas as pd

image_path = "../Resources/TrainingImages/"

df = pd.read_csv("../Resources/train.csv")

df["file_name"] = df["file_name"].str.replace("train_data/", "")
```

```
df.head()
```

Unnamed: 0		file_name	label
0	0	a6dcb93f596a43249135678dfcfc17ea.jpg	1
1	1	041be3153810433ab146bc97d5af505c.jpg	0
2	2	615df26ce9494e5db2f70e57ce7a3a4f.jpg	1
3	3	8542fe161d9147be8e835e50c0de39cd.jpg	0
4	4	5d81fa12bc3b4cea8c94a6700a477cf2.jpg	1

```
for index, row in df.iterrows():
    os.rename(f'{image_path}{row["file_name"]}', f'{image_path}{row["label"]}/{row["file_name"]}')
```

```
#Randomly sample 20% of image set
import random

test_sample = int(len(os.listdir(f'{image_path}/0')) * 0.2)
test_files = random.sample(os.listdir(f'{image_path}/0'), test_sample)
```

```
test_files[:5]
```

```
['4d7f187ea7104899b0bf20b5219c69d1.jpg',
 'a6faa0b9b90440298b1caa88f610efad.jpg',
 'eb3dd55bcfcb4c4dac8742b0004b4a30.jpg',
 '2b6211e5e47d43d3a3b7ee0c06182a0a.jpg',
 'b2255267b06b448984d3bb1695036c03.jpg']
```

```
#Move real images to testing
for row in test_files:
    os.rename(f'{image_path}0/{row}', f'../Resources/TestingImages/0/{row}')
```

```
#Move AI images to testing
test_files2 = random.sample(os.listdir(f'{image_path}/1'), test_sample)
```

```
for row in test_files2:
    os.rename(f'{image_path}1/{row}', f'../Resources/TestingImages/1/{row}')
```

Splitting Testing & Training Data

Classes:

0 = Human-created

1 = AI-generated

With the file structure we created, we can use

`labels="inferred"`

This classifies images based explicitly on their directory location

```
#Split testing and training data
train_image_path = "./Resources/TrainingImages/"
test_image_path = "./Resources/TestingImages/"
image_size = (255, 255)

y = np.array(df["label"])
X_train = tf.keras.preprocessing.image_dataset_from_directory(
    train_image_path,
    labels="inferred",
    image_size=image_size,
    interpolation="bilinear"
)
X_test = tf.keras.preprocessing.image_dataset_from_directory(
    test_image_path,
    labels="inferred",
    image_size=image_size,
    interpolation="bilinear"
)
```

Found 63960 files belonging to 2 classes.

Found 15990 files belonging to 2 classes.

Creating the Learning Model: Convolutional Neural Network

Convolutional and
Max Pooling layers
effectively halve the
output shape

```
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 253, 253, 32)	896
max_pooling2d_5 (MaxPooling2D)	(None, 126, 126, 32)	0
conv2d_6 (Conv2D)	(None, 124, 124, 64)	18,496
max_pooling2d_6 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_7 (Conv2D)	(None, 60, 60, 64)	36,928
max_pooling2d_7 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_8 (Conv2D)	(None, 28, 28, 64)	36,928
max_pooling2d_8 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_9 (Conv2D)	(None, 12, 12, 64)	36,928
flatten_1 (Flatten)	(None, 9216)	0
dense_2 (Dense)	(None, 64)	589,888
dense_3 (Dense)	(None, 2)	130

Total params: 720,194 (2.75 MB)
Trainable params: 720,194 (2.75 MB)
Non-trainable params: 0 (0.00 B)

Testing the Learning Model: Accuracy & Loss

Each image
is scaled to
255 by 255
pixels, with
3 layers for
red, green,
and blue

```
model.compile(optimizer='adam',  
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              metrics=['accuracy'])
```

```
history = model.fit(X_train, epochs=10)
```

```
Epoch 1/10  
1999/1999 ————— 1103s 550ms/step - accuracy: 0.8221 - loss: 0.7447  
Epoch 2/10  
1999/1999 ————— 1061s 531ms/step - accuracy: 0.9176 - loss: 0.2154  
Epoch 3/10  
1999/1999 ————— 1053s 527ms/step - accuracy: 0.9293 - loss: 0.1904  
Epoch 4/10  
1999/1999 ————— 1051s 526ms/step - accuracy: 0.9389 - loss: 0.1660  
Epoch 5/10  
1999/1999 ————— 1054s 527ms/step - accuracy: 0.9418 - loss: 0.1552  
Epoch 6/10  
1999/1999 ————— 1052s 526ms/step - accuracy: 0.9454 - loss: 0.1423  
Epoch 7/10  
1999/1999 ————— 1051s 526ms/step - accuracy: 0.9495 - loss: 0.1400  
Epoch 8/10  
1999/1999 ————— 1052s 526ms/step - accuracy: 0.9557 - loss: 0.1161  
Epoch 9/10  
1999/1999 ————— 1052s 526ms/step - accuracy: 0.9576 - loss: 0.1203  
Epoch 10/10  
1999/1999 ————— 1052s 526ms/step - accuracy: 0.9619 - loss: 0.1026
```

This means
that each
image file has
AT LEAST
195,075 points
of data to
process!

Testing Images from URL

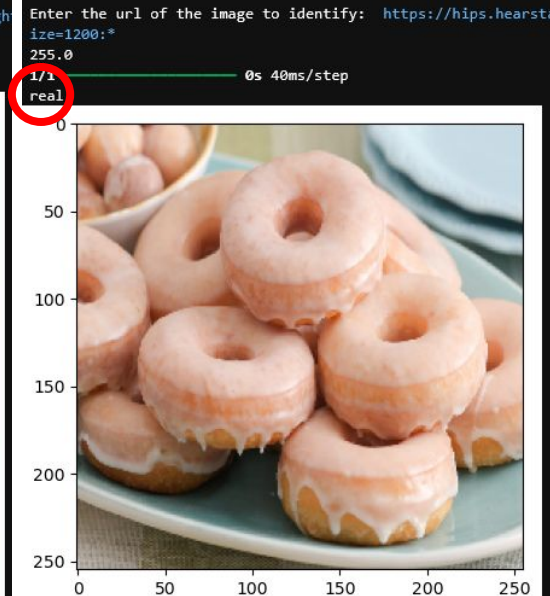

```
# Test images from url
import urllib.request
import io

image_size = (255, 255)
label = ["real", "ai"]

def predictImage(url):
    with urllib.request.urlopen(url) as rawdata:
        img = tf.keras.preprocessing.image.load_img(io.BytesIO(rawdata.read()), target_size=image_size)
        x = tf.keras.preprocessing.image.img_to_array(img)
        x = x[np.newaxis, :, :, :]
        print(np.max(x))
        plt.imshow(img)
        predictions = model.predict(x)
        prediction = np.argmax(predictions)
        print(label[prediction])

url = input("Enter the url of the image to identify: ")
predictImage(url)

Enter the url of the image to identify: https://media.licdn.com/dms/image/v2/D5622AQEsa052pRM_A/feedshare-shrink_800/B547483647&v=beta&t=EIH39mnR31fdC9HNOgxoof12RcZb71Hwkb_GY9-Eji
```




Testing Images from URL



They're not all winners, though...

Conclusion:
Training image dataset
severely lacking in
chinchillas




NEXT STEPS:

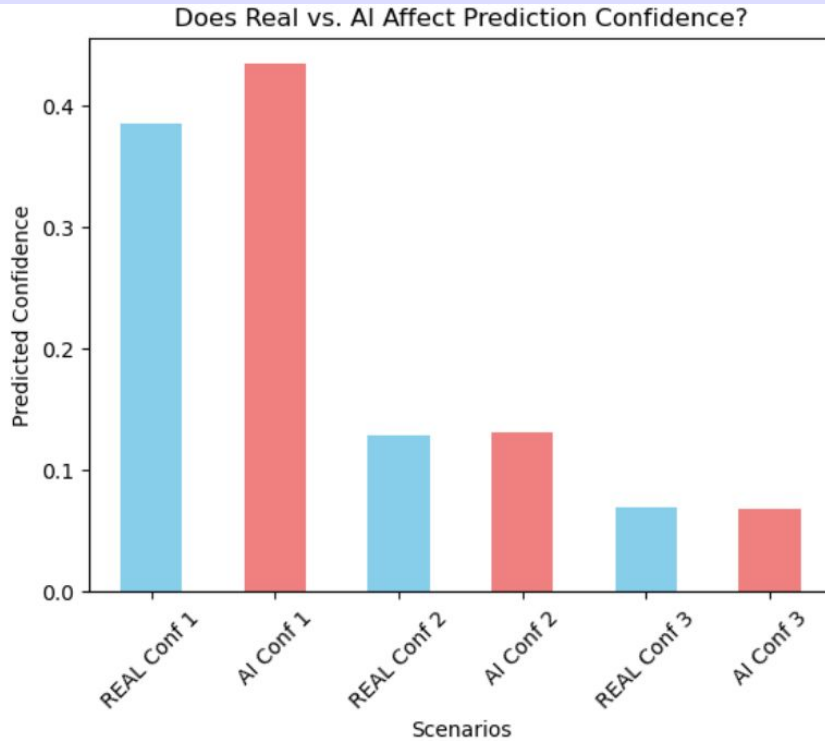
Due to timing constraints, we were unable to analyze output from our real vs. AI prediction model. However, it did make sense to build methodology for analyzing in the future.

- Used a manageable subset of training images (2000 random)
- Used class example of image prediction model to group images vs. manual categorization.
- Build a scatter-plot based upon prediction confidence to see if we could “visually” spot trends.

Analyzing the Training Imagery



Did the image categorization engine (VGG19) have more or less confidence in real vs. AI images?



Scenarios		Predicted Confidence
0	REAL Conf 1	0.385334
1	AI Conf 1	0.434599
2	REAL Conf 2	0.129261
3	AI Conf 2	0.131992
4	REAL Conf 3	0.070293
5	AI Conf 3	0.068514

1st Prediction,
mean confidence

Why so many lab_coat(s)?

```
# Look up Real unique values
```

```
real_uniques = df_merged[df_merged['label'] == 0]['predict_1'].value_counts()  
real_uniques.head()
```

```
predict_1  
lab_coat    34  
plate       23  
jean        22  
bikini      15  
alp         14  
Name: count, dtype: int64
```



These unique values were
classified the most out of 1,021
real images.

```
# Look up AI unique values
```

```
ai_uniques = df_merged[df_merged['label'] == 1]['predict_1'].value_counts()  
ai_uniques.head()
```

```
predict_1  
lab_coat    28  
rapeseed    21  
wig          19  
cucumber    14  
maillot     14  
Name: count, dtype: int64
```



These unique values were
classified the most out of 979
AI generated images.

Real Lab Coat Imagery vs. Prediction Confidence



AI Lab Coat Imagery vs. Prediction Confidence




What do plate(s) look like?

```
# Look up Real unique values
```

```
real_uniques = df_merged[df_merged['label'] == 0]['predict_1'].value_counts()  
real_uniques.head()
```


```
predict_1  
lab_coat    34  
plate       23  
jean        22  
bikini      15  
alp         14  
Name: count, dtype: int64
```




```
# Look up AI counts for plate
```

```
ai_plate_uniques = df_merged[(df_merged['label'] == 1) & (df_merged['predict_1'] == "plate")]['predict_1'].value_counts()  
ai_plate_uniques
```

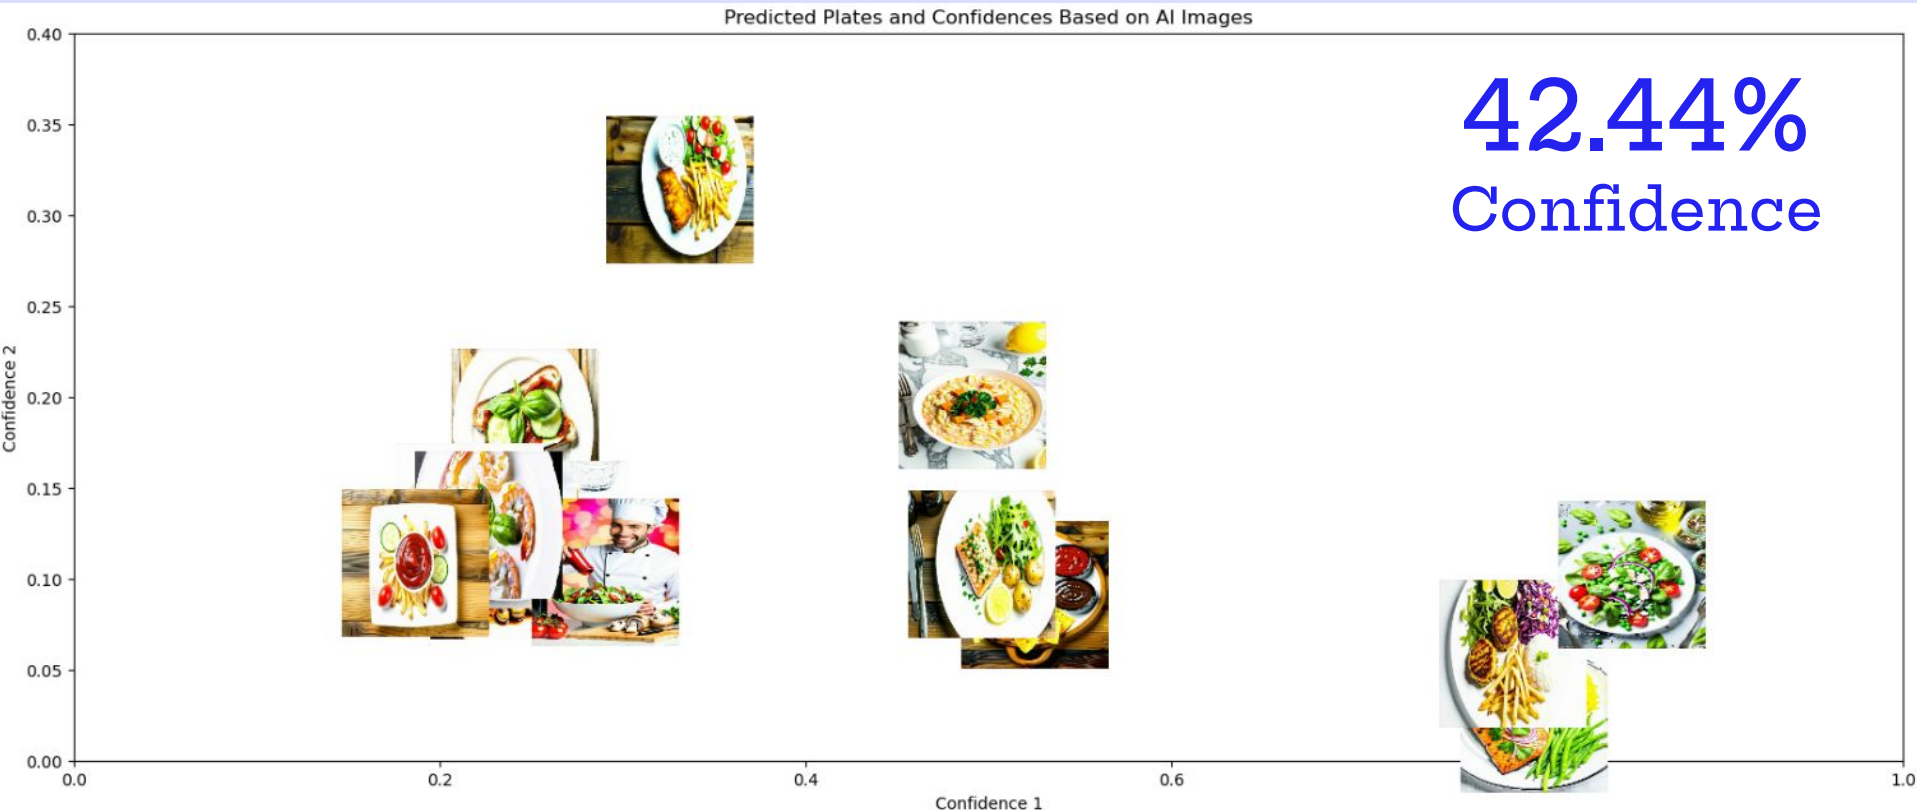
```
predict_1  
plate      14  
Name: count, dtype: int64
```



Real Plate Imagery vs. Prediction Confidence



AI Plate Imagery vs. Prediction Confidence




Finally, what is rapeseed and why did AI over-index with it?

```
# Look up real counts for rapeseed
```

```
real_rapeseed_uniques = real_rapeseed[(real_rapeseed['label'] == 0) & (real_rapeseed['predict_1'] == "rapeseed")]['predict_1'].value_counts()  
real_rapeseed_uniques
```

```
predict_1  
rapeseed    6  
Name: count, dtype: int64
```



```
# Look up AI unique values
```

```
ai_uniques = df_merged[df_merged['label'] == 1]['predict_1'].value_counts()  
ai_uniques.head()
```

```
predict_1  
lab_coat    28  
rapeseed    21  
wig         19  
cucumber    14  
maillot     14  
Name: count, dtype: int64
```



what is rapeseed

All Images Videos Shopping Short videos Forums Web More

Rapeseed

Plant :

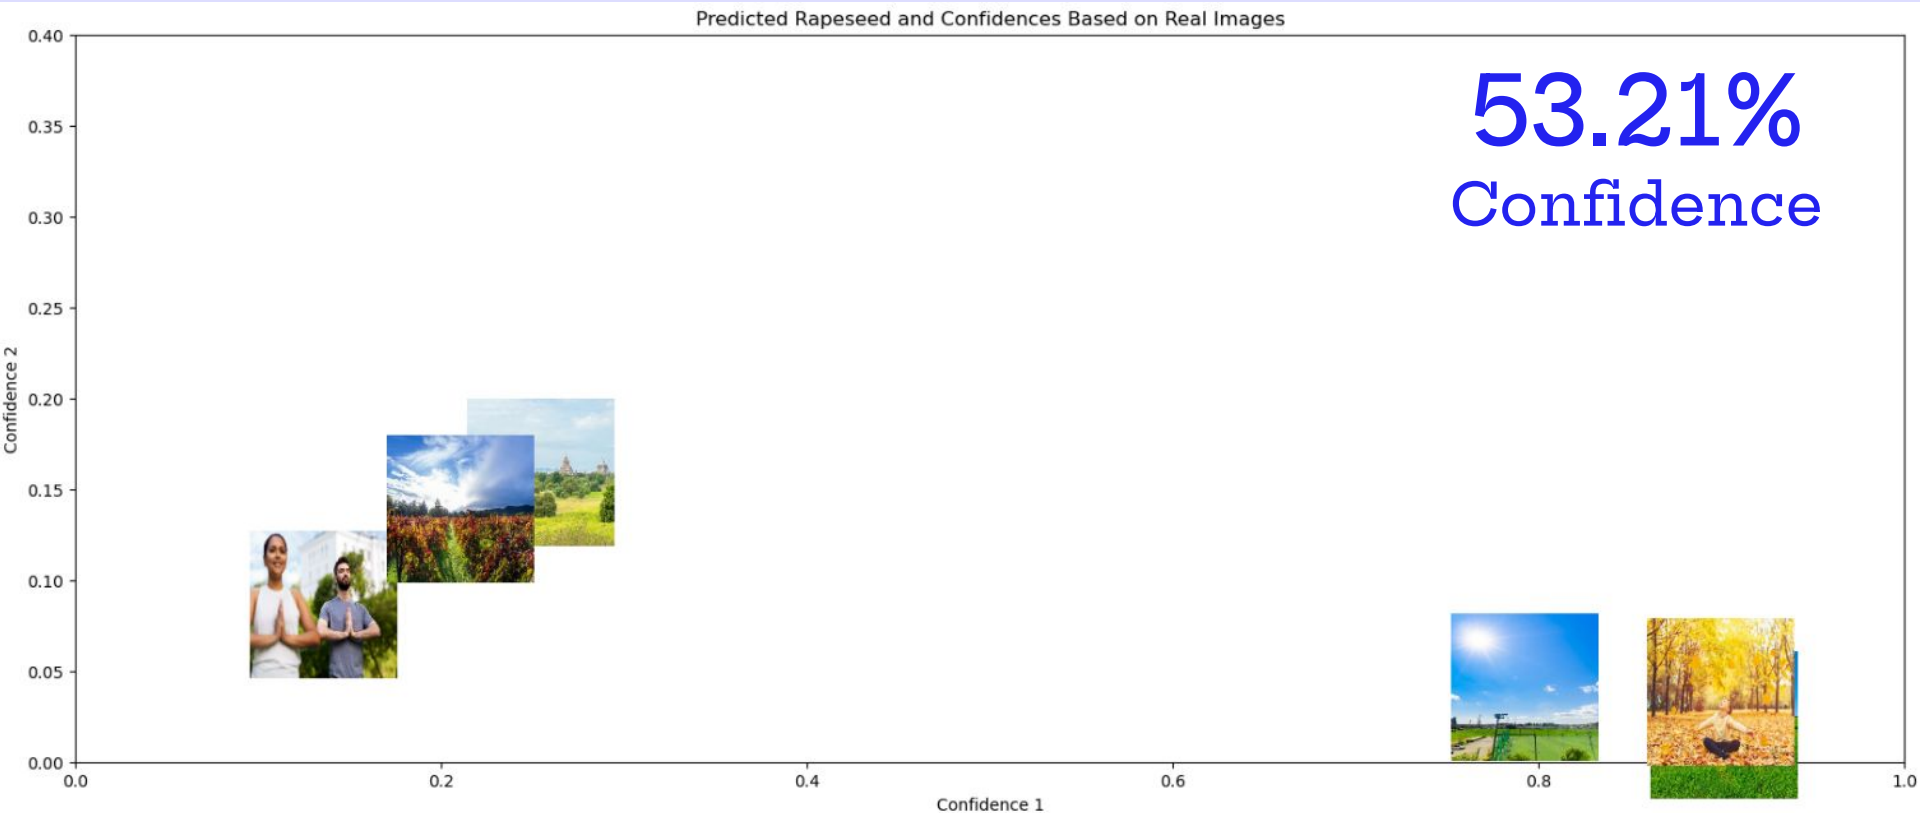
Overview

Benefits

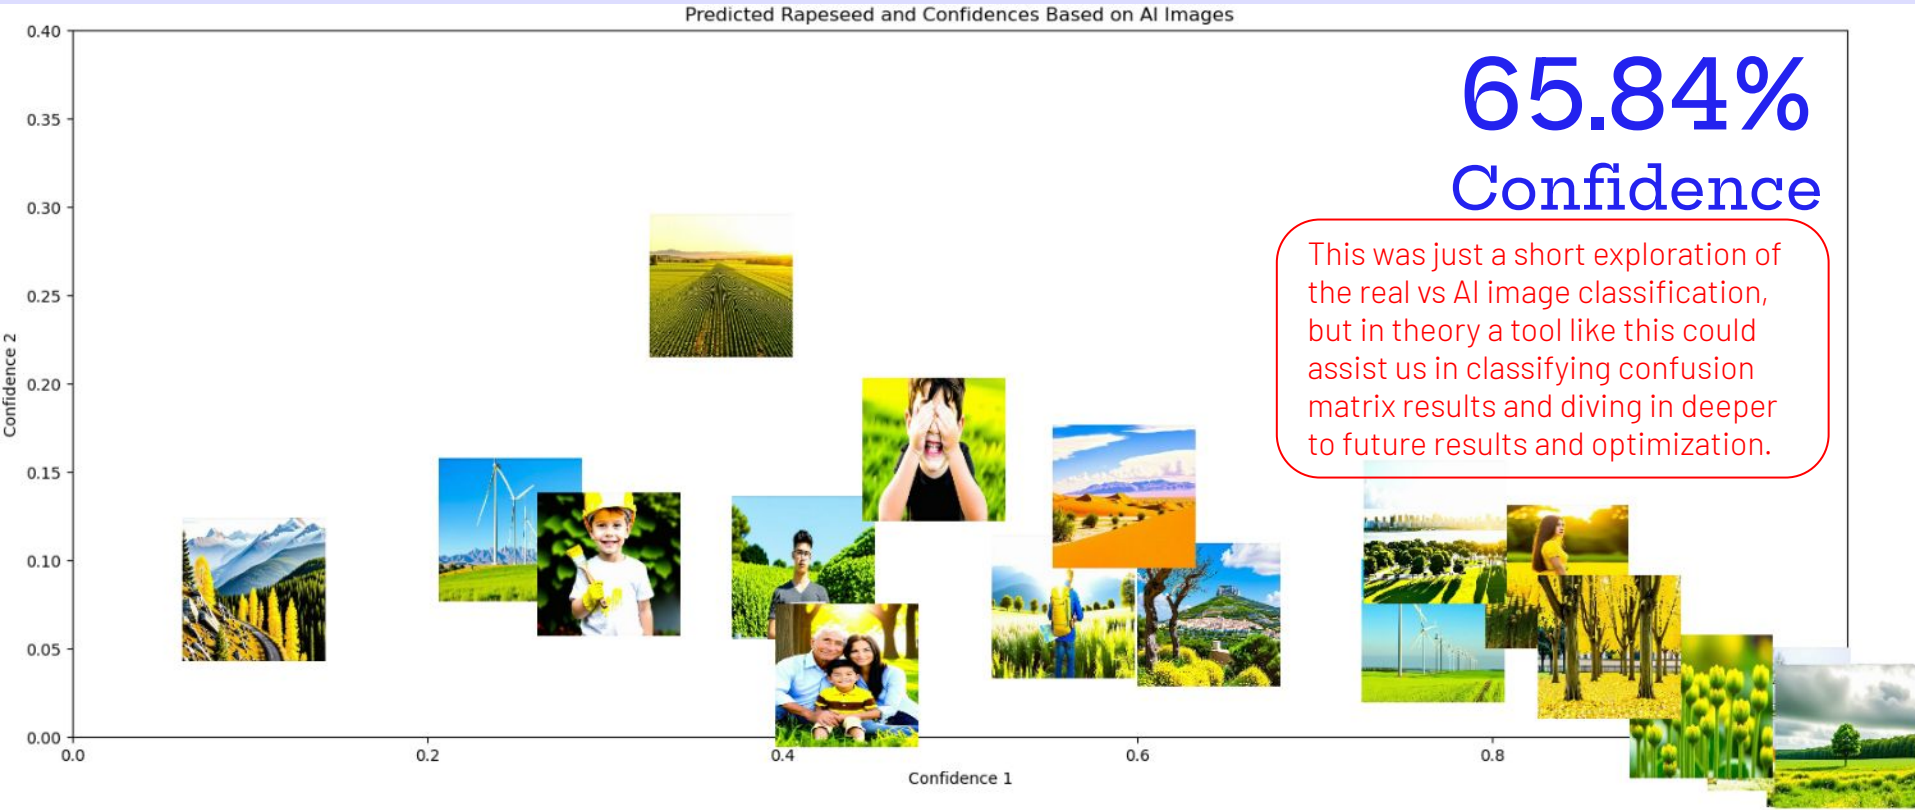
Lower classifications



Real Rapeseed Imagery vs. Prediction Confidence



AI Rapeseed Imagery vs. Prediction Confidence



FLASK-APP BUILD

26

We wanted to be able to create a webpage that would be able to predict uploaded images based on the train model created on whether an image was AI-generated or Human-generated.

The app.py was built so it would be able to run either locally with python for single uploads or deployed with a cloud platform service (heroku) that would be able to support multiple uploads at the same time.

```
py 6 X
ers > ToniMak > Desktop > pretrained_testing(PROJECT4) > flask_app > app.py > predict

# Load environment variables
load_dotenv()

# AWS credentials
AWS_ACCESS_KEY_ID = os.getenv("AWS_ACCESS_KEY_ID")
AWS_SECRET_ACCESS_KEY = os.getenv("AWS_SECRET_ACCESS_KEY")
S3_BUCKET_NAME = os.getenv("S3_BUCKET_NAME")
MODEL_PATH = "models/model_05.h5"
LOCAL_MODEL_PATH = "/tmp/model_05.h5"

# Set LOCAL_MODEL_PATH based on OS
if os.name == 'nt': # Windows
    # Create 'tmp' folder in your current working directory if not exists
    LOCAL_MODEL_PATH = os.path.join(os.getcwd(), "tmp", "model_05.h5")
    os.makedirs(os.path.dirname(LOCAL_MODEL_PATH), exist_ok=True)
else:
    # On Heroku, use the absolute path in /tmp
    LOCAL_MODEL_PATH = "/tmp/model_05.h5"

# Initialize S3 client
s3 = boto3.client(
    "s3",
    aws_access_key_id=AWS_ACCESS_KEY_ID,
    aws_secret_access_key=AWS_SECRET_ACCESS_KEY
)
```

```
def predict():
    if file.filename == "":
        return jsonify({"error": "No selected file"}), 400

    try:
        # Convert the uploaded file to a BytesIO object
        img_bytes = BytesIO(file.read())

        # Load and preprocess the image
        img = image.load_img(img_bytes, target_size=(255, 255))
        img_array = image.img_to_array(img)
        img_array = np.expand_dims(img_array, axis=0) / 255

        # Make prediction (get raw output from the model)
        preds = model.predict(img_array)

        # Debug: log raw prediction values
        print("Raw prediction:", preds)

        # Check output shape to decide which branch to use
        if preds.shape[-1] == 1:
            # Model with a single output neuron
            raw_value = preds[0][0]
            # Apply sigmoid manually to convert to probability
            prob = tf.nn.sigmoid(raw_value).numpy()
            if prob >= 0.5:
                label = "Human-generated"
                confidence = prob * 100
            else:
                label = "AI-generated"
                confidence = (1 - prob) * 100
```

INDEX.HTML BUILD

27

```
index.html X
ers > TonMak > Desktop > pretrained_testing(PROJECT4) > flask_app > templates > index.html > head > style > #warning
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>AI-Generated vs Human-Generated Image Classifier</title>
  <style>
    body {
      font-family: Arial, sans-serif;
      text-align: center;
      margin: 50px;
    }
    #warning {
      color: red;
      font-weight: bold;
    }
  </style>
</head>
<body>
  <h1>Upload an Image to Classify</h1>

  <!-- User Guidance Message -->
  <p id="warning">⚠ Please upload a <strong>square-shaped image</strong> to minimize distortion when resizing.</p>

  <form id="uploadForm" enctype="multipart/form-data">
    <label for="file">Select an image:</label>
    <input type="file" name="file" accept="image/*" required><br><br>
    <input type="submit" value="Upload and Predict">
  </form>

  <div id="predictionResult"></div>
</body>
</html>
```

```
<script>
  document.getElementById('uploadForm').addEventListener('submit', async (e) => {
    e.preventDefault();

    let form = document.getElementById('uploadForm');
    let formData = new FormData(form);
    document.getElementById('predictionResult').innerHTML = "<p>Processing... Please wait.</p>";

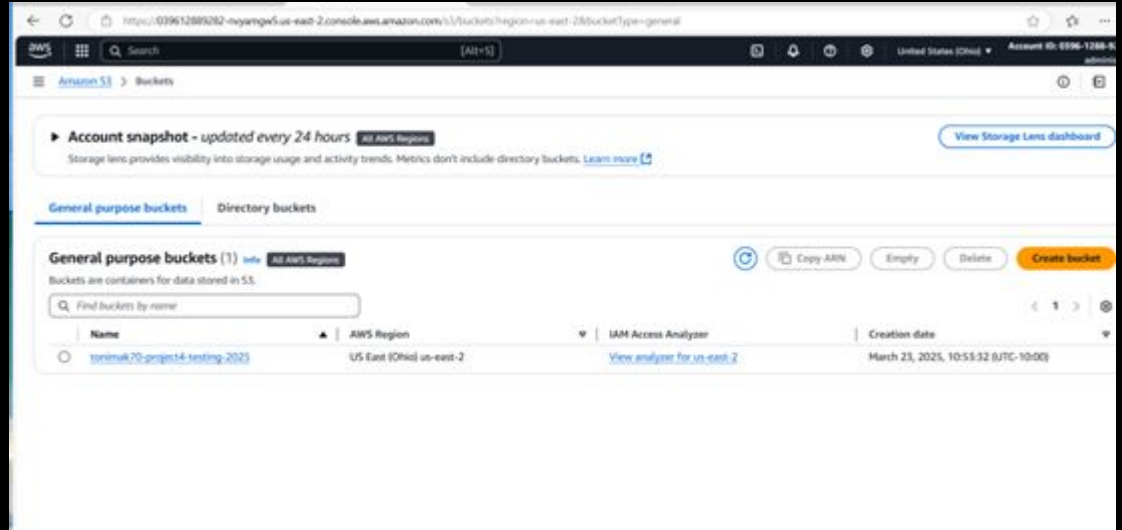
    let baseUrl = window.location.origin; // sets to localhost or Heroku

    try {
      let response = await fetch(`${baseUrl}/predict`, {
        method: 'POST',
        body: formData
      });

      let result = await response.json();
      let predictionDiv = document.getElementById('predictionResult');

      if (result.predictions) {
        let output = "<h2>Prediction Result</h2><ul>";
        result.predictions.forEach((pred, index) => {
          output += "<li><strong>${index + 1}. ${pred.label}</strong> - Confidence: ${pred.confidence.toFixed(2)}%</li>";
        });
        output += "</ul>";
        predictionDiv.innerHTML = output;
      } else {
        predictionDiv.innerHTML = "<p>Error: ${result.error}</p>";
      }
    } catch (error) {
      console.error(error);
      document.getElementById('predictionResult').innerHTML = "<p>Failed to process image. Please try again.</p>";
    }
  });
</script>
</body>
</html>
```

CLOUD STORAGE SERVICES




Used AWS S3 cloud storage used to store and retrieve our large data files for the heroku deployment

CLOUD PLATFORM SERVICES

<https://app-flask-xception-7d7b178928d0.herokuapp.com/>

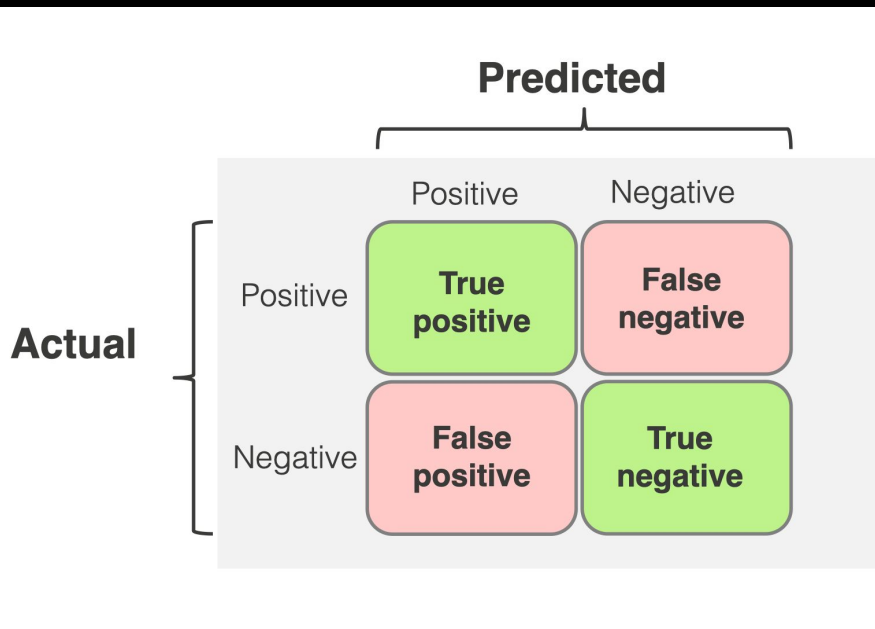


Used HEROKU, a cloud platform, to deploy our application with the intent to have multiple users upload images at the same time.



But wait, what would we do if we had more time to investigate?

30



Study the results more closely

- **Confusion matrix of images**
- **Shared characteristics across False Positives and False Negatives**
- **Build visualisations to further understand the model's accuracy in relation to real vs. AI images.**

Questions?

THANK YOU